



PREMISE:
Three tiny runs support costly large-scale decisions.



Goal:
audit whether these tiny runs can be trusted to guide large decisions.




1 CHECKPOINT-NOISE PROXY

CAN ONE RUN CERTIFY CORRELATION?



Claim: a single run's last checkpoints can substitute for rerunning seeds - advertised correlation 0.90

HOW WE AUDIT

one run → many checkpoints



CEILING LOGIC



COMPATIBLE WITH ITS CEILING BUT UNCERTIFIABLE

2 SOURCE EQUIVALENCE

CAN INIT-ONLY NOISE STAND IN?

Claim: an init-only noise floor can stand in for the all-source seed-noise floor.

HOW WE AUDIT

initialization (I)
(different seeds)

data order (D)
(different orders)

init-only arm (I)
many seeds,
fixed order

order-only arm (D)
many orders,
fixed init

σ_{init} = SD over
init-only arm
 σ_{order} = SD over
order-only arm


$r = \sigma_{init} / \sigma_{order}$



undecidable if
no CI fits

why single-source is not enough

Full model: $f(I, D) = \mu + a(I) + b(D) + c(I, D)$
Single-source arms estimate $a(I)$ or $b(D)$ only.
Interaction $c(I, D)$ is not identified.






UNDECIDABLE

3 DECISION-INSTRUMENT SNR

DO SMALL-SCALE RANKINGS CARRY SIGNAL?

Claim: small-scale rankings carry enough signal over seed noise to support decisions.

HOW WE AUDIT




- recipe mean
- I within-recipe SD(seed noise)

signal strength among recipes



flat up to 60M, jumps ×2.1 at 60M → 90M.



NOISE-DOMINATED (≤60M, accuracy readout)